


B+ Trees Indexing

Dr. Shaheen Khatoon



Indexes: Review

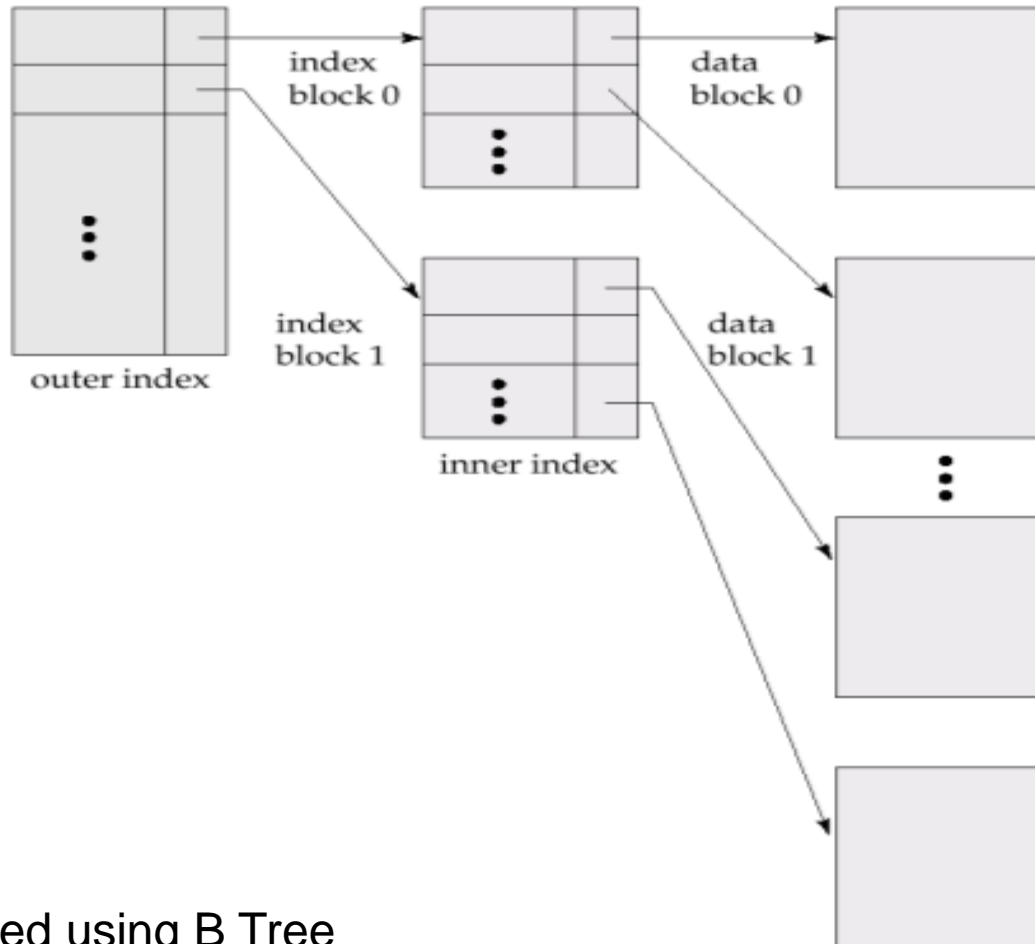
- ❑ Indexing mechanisms used to speed up access to desired data.
- ❑ **Search Key** - attribute to set of attributes used to look up records in a file.
- ❑ An **index file** consists of records (called index entries) of the form

search-key	pointer
------------	---------
- ❑ Index files are typically much smaller than the original file
- ❑ Two basic kinds of implementations:
 - **Tree indices**: search keys are stored in sorted order!
 - **Hash indices**: search keys are distributed uniformly across “buckets” using a “hash function”.

Multilevel Index: review

- ❑ If primary index does not fit in memory, access becomes expensive.
- ❑ To reduce number of disk accesses to index records, treat primary index kept on disk as a sequential file and construct a sparse index on it.
- ❑ **outer index** – a sparse index of primary index
- ❑ **inner index** – the primary index file
- ❑ If even outer index is too large to fit in main memory, yet another level of index can be created, and so on.
- ❑ Indices at all levels must be **updated on insertion or deletion from the file**.

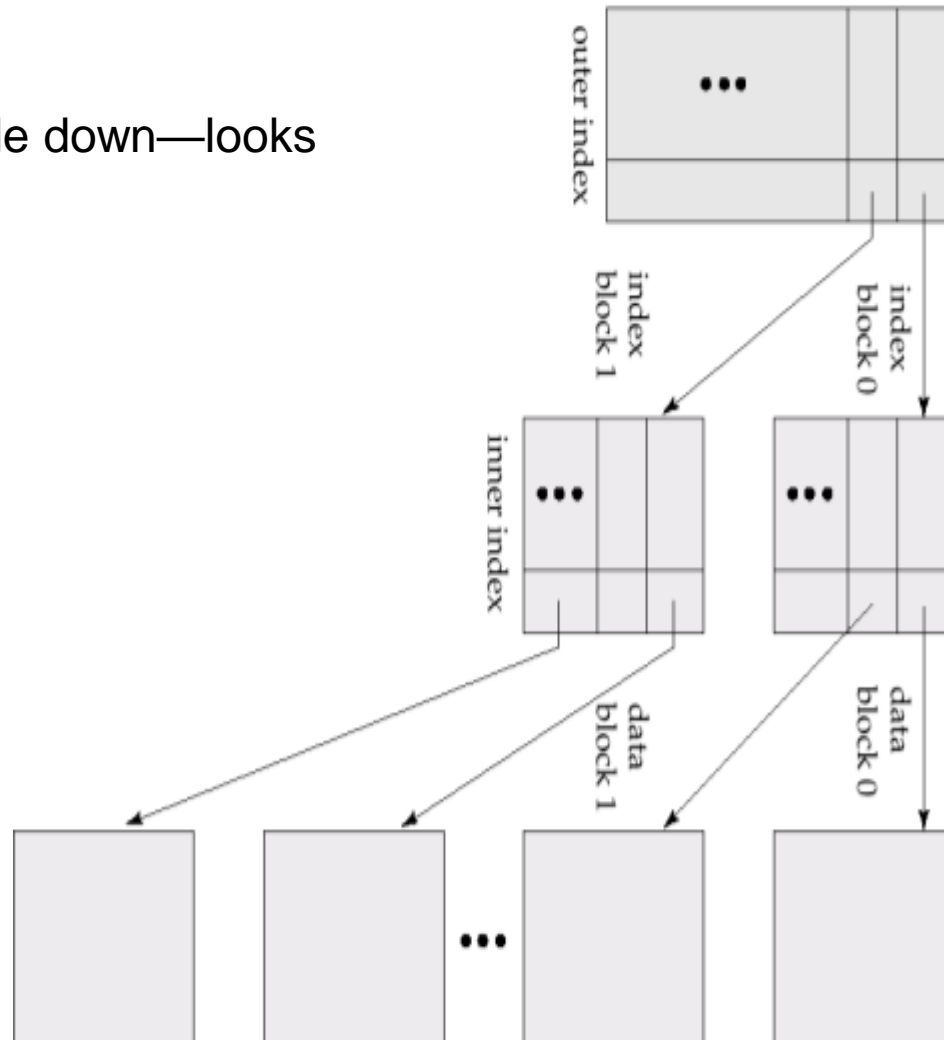
Multilevel Index (Cont.)



Implemented using B Tree


Multilevel Index (Cont.)

Turn multi-lev index upside down—looks like a tree




What is a binary tree?

- ❑ *Property 1:* each node can have up to two successor nodes.
- ❑ *Property 2:* a unique path exists from the root to every other node




Not a valid binary tree




Some terminology

- ❑ The successor nodes of a node are called its *children*
- ❑ The predecessor node of a node is called its *parent*
- ❑ The "beginning" node is called the *root* (has no parent)
- ❑ A node without *children* is called a *leaf*



Some terminology (cont'd)

- ❑ Nodes are organized in levels (indexed from 0).
- ❑ **Height of a tree h :** $\#levels = L$
- ❑ **Full tree:** every node has exactly two children *and* all the leaves are on the same level.




B + tree Indexing: Concept

- ❑ Each node can have up to **n** children
 - where n is fan out (FO), maximum degree or order of tree.
 - Each node can store n pointer and (n-1) search values
- ❑ Allow efficient and fast exploration at the expense of using slightly more space.
- ❑ Support more efficiently queries like:
`SELECT * FROM R WHERE a = 11`
- ❑ `SELECT * FROM R WHERE 0 ≤ b and b < 42`

B+ Tree node Structure


□ Typical node Structure



- K_i are the search-key values
- P_i are pointers to children (for non-leaf nodes) or pointers to records (for leaf nodes).
- The search-keys in a node are ordered $K_1 < K_2 < K_3 < \dots < K_{n-1}$
- Each node stored in **one disk block**

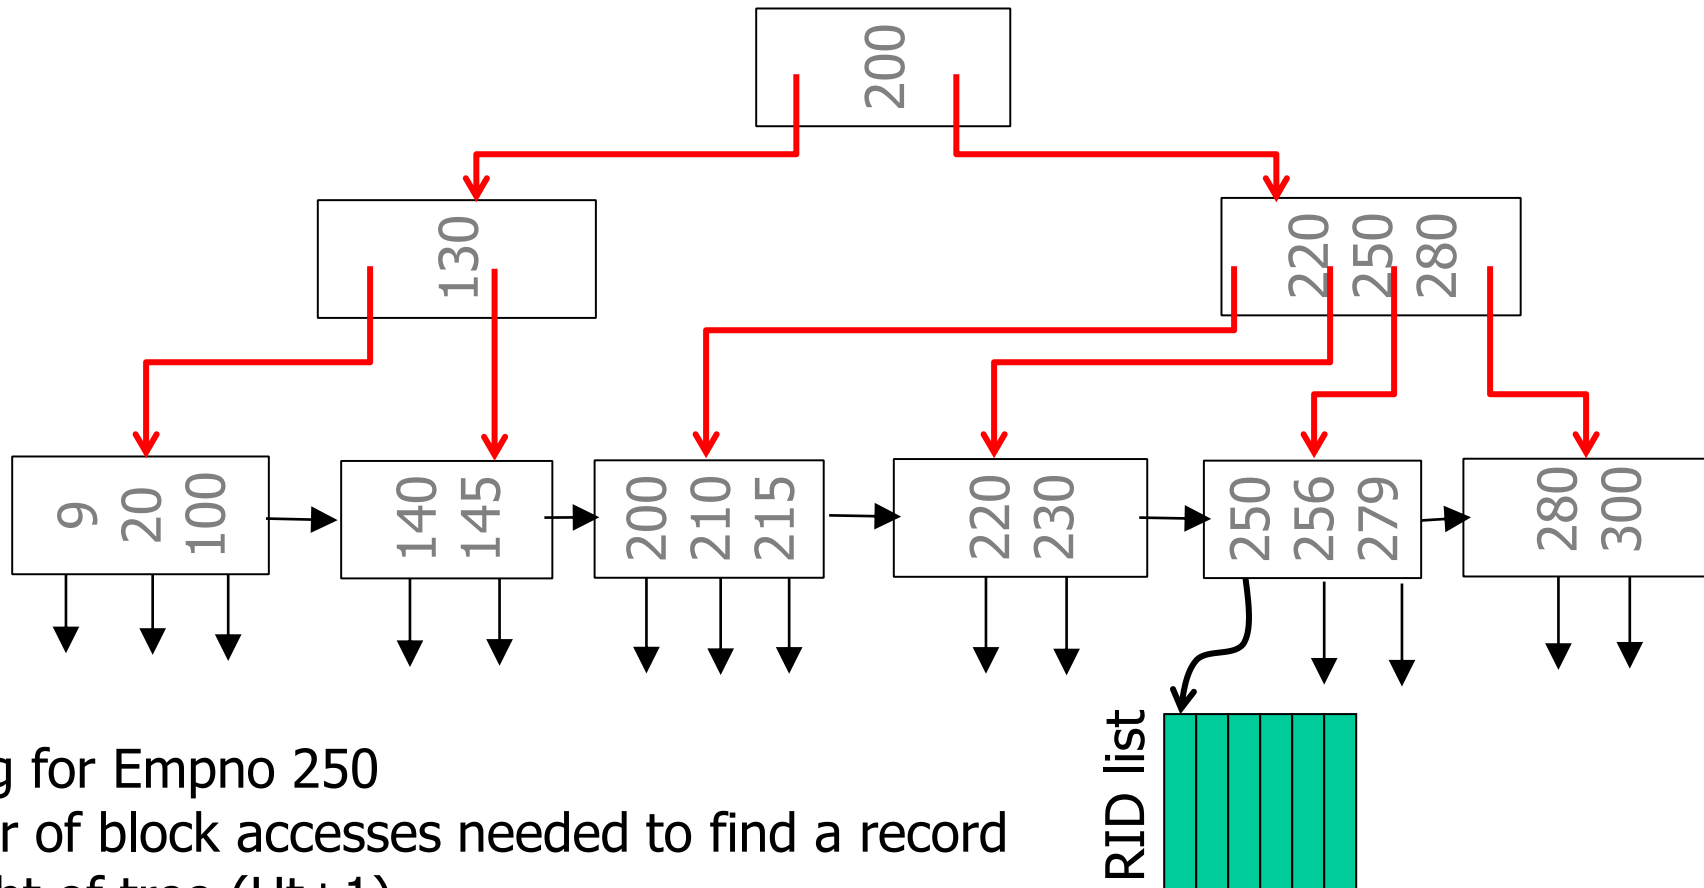
Leaf Node in B+ Tree

- For $i = 1, 2, \dots, n-1$, pointer P_i points to a file record with search-key value K_i
- If L_i, L_j are leaf nodes and $i < j$, L_i 's search-key values are less than L_j 's search-key values
 - P_n points to next leaf node in search-key order



B-tree Indexing: Example

- ❑ B+ tree of order 4
 - Maximum 4 pointer and 3 keys



Looking for Empno 250
Number of block accesses needed to find a record
= Height of tree (Ht+1)
In this case $3+1 = 4$

Maximum Capacity of B+ Tree

- ❑ In a 3 level B+ tree with 100 key and 101 pointers in each node:
 - we can index as many as a million records ($101 * 101 * 100 = 1$ million+).
- ❑ With such an index, accessing a database record requires only 3 index-block read and 1 data-block read. *Very efficient!*
- ❑ Given the depth of the tree is 4, with 3 pointers
 - Index $3 \times 3 \times 3 \times 2 = 54$ records in database. This is the maximum capacity of this B+ tree index.
- ❑ If we use a B+ tree with **order, p = 100**, maximum how many records or blocks can we index using a **3-level** tree?
 - $100 \times 100 \times 99 = 990,000$ (almost a million)

Searching in B+ Tree

□ Generalization of binary search

1. Give a search key, start from root node
2. If key is present in root node then success, else
3. If current node is a leaf node and key not present in node, then key not in database
4. Search for tree pointer p_i such that $k_{i-1} < k_i < k_{i+1}$
5. Return step 2 to continue search


Inserting a Data Entry into a B+ Tree

- ❑ Find correct leaf L .
- ❑ Put data entry onto L .
 - If L has enough space, *done!*
 - Else, must split L (into L and a new node $L2$)
 - Redistribute entries evenly, copy up middle key.
 - Insert index entry pointing to $L2$ into parent of L .
- ❑ This can happen recursively
 - To split index node, redistribute entries evenly, but push up middle key. (Contrast with leaf splits.)
- ❑ Splits “grow” tree; root split increases height.
 - Tree growth: gets wider or one level taller at top.

Example of B+ tree

This table shows a B+ tree. As the example illustrates this tree does not have a full index page. (We have room for one more key and pointer in the root page.) In addition, one of the data pages contains empty slots

B+ Tree with four keys



Illustrations of the insert algorithm



- ❑ The key value determines a record's placement in a B+ tree.
- ❑ The **leaf pages** are maintained in **sequential order**
- ❑ **Doubly linked** list connects each leaf page with its **sibling page(s)**.
- ❑ This doubly linked list **speeds data movement** as the pages grow and contract.

Illustrations of the insert algorithm

- ❑ The following examples illustrate each of the **insert** scenarios.
- ❑ We begin with the simplest scenario:
 - inserting a record into a leaf page that is not full.
 - Since only the leaf node containing **25 and 30** contains expansion room, we're going to insert a record with a key value of **28** into the B+ tree.
 - The following figures shows the result of this addition.

Add to a non-full tree

- Add Record with Key 28



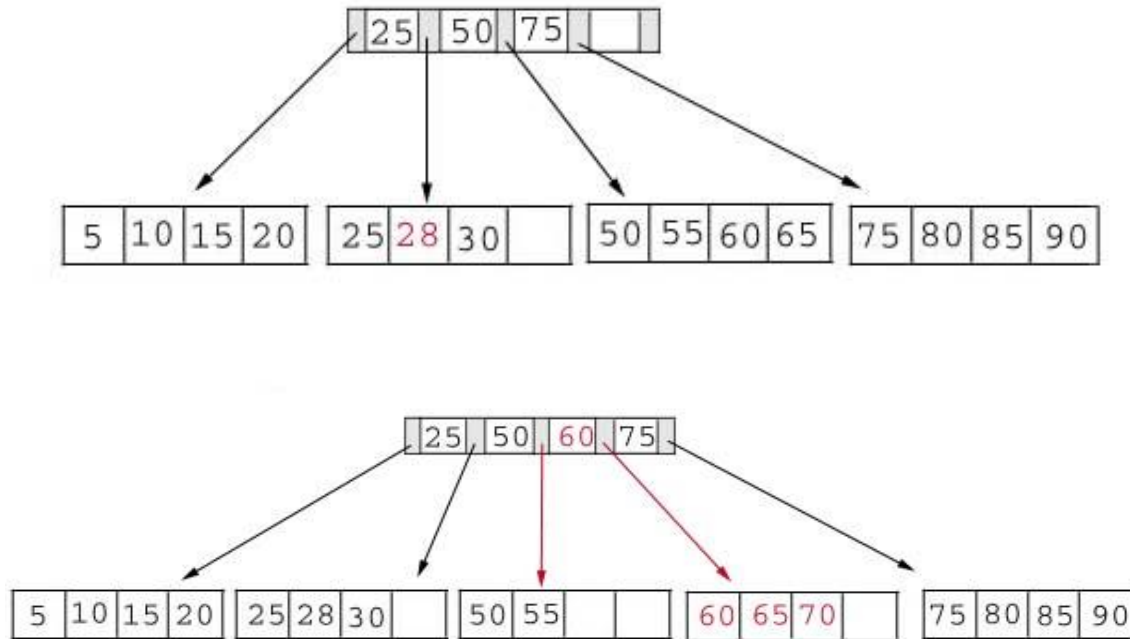
Adding a record when the **leaf page is full** but the **index page is not**

- ❑ we're going to insert a record with a key value of **70** into our B+ tree.
- ❑ This record should go in the leaf page containing 50, 55, 60, and 65.
- ❑ Unfortunately this **page is full**. This means that we must split the page as follows

❑ Left Leaf Page	Right Leaf Page
50 55	60 65 70


Add Record with Key 70 : Leaf node split example

- ❑ The middle key of 60 is placed in the index page between 50 and 75.
- ❑ The following table shows the B+ tree after the addition of 70.




Adding a record when both the **leaf page** and the **index** page are **full**

- ❑ Add **95** to B+ tree.
- ❑ This record belongs in the page containing **75, 80, 85, and 90**. Since this page is full, split it into two pages:
- ❑ **Left Leaf Page Right Leaf Page**
- ❑ **75 80 85 90 95**

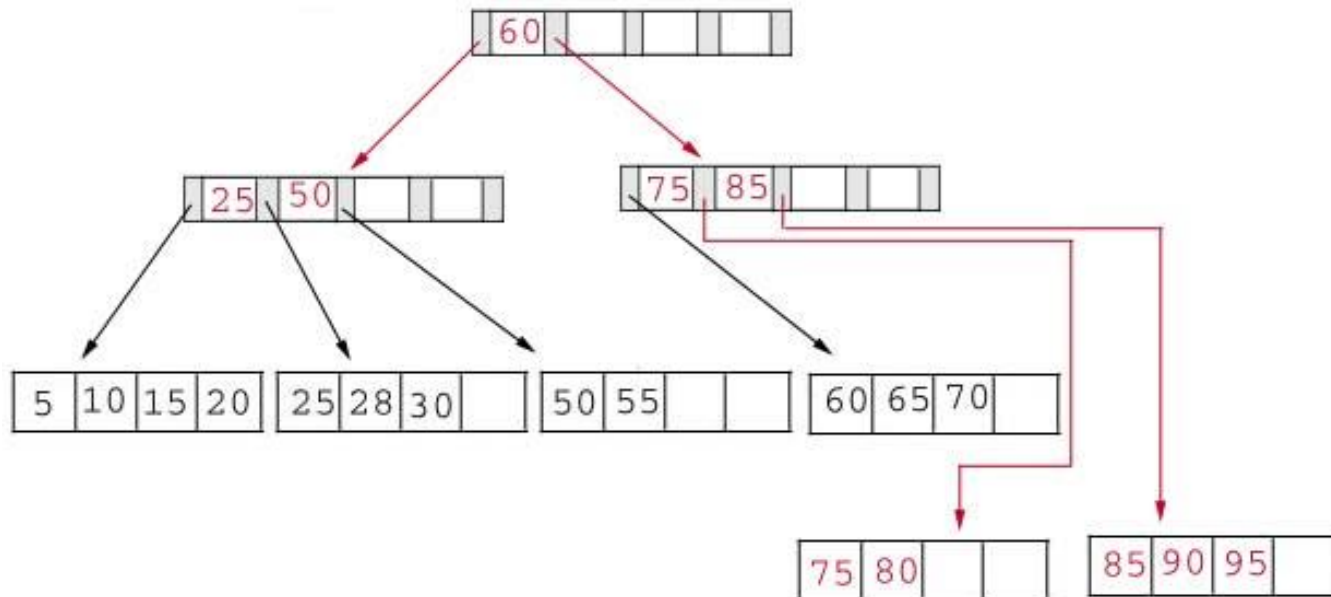


-
- ❑ The middle key, **85**, rises to the index page.
 - ❑ Unfortunately, the index page is also full, so we split the index page:

Left Index Page	Right Index Page	New Index Page
25 50	75 85	60



The following table illustrates the addition of the record containing **95** to the B+ tree.

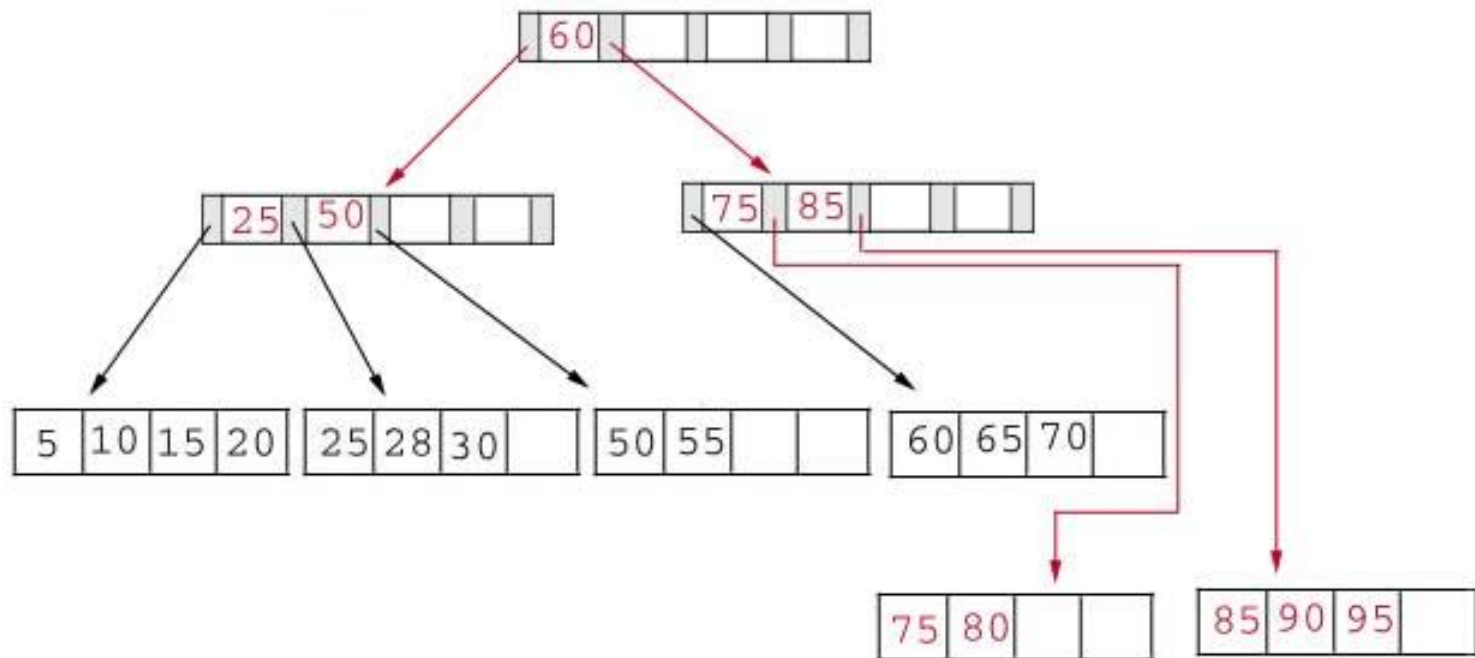


Deleting a Data Entry from a B+ Tree

- ❑ Start at root, find leaf L where entry belongs.
- ❑ Remove the entry.
 - If L is at least half-full, *done!*
 - If L has less than half entries,
 - Try to **re-distribute**, borrowing from sibling (*adjacent node with same parent as L*).
 - If re-distribution fails, merge L and sibling.
- ❑ If merge occurred, must delete entry (pointing to L or sibling) from parent of L .
- ❑ Merge could propagate to root, decreasing height.

Deleting Keys from a B+ tree


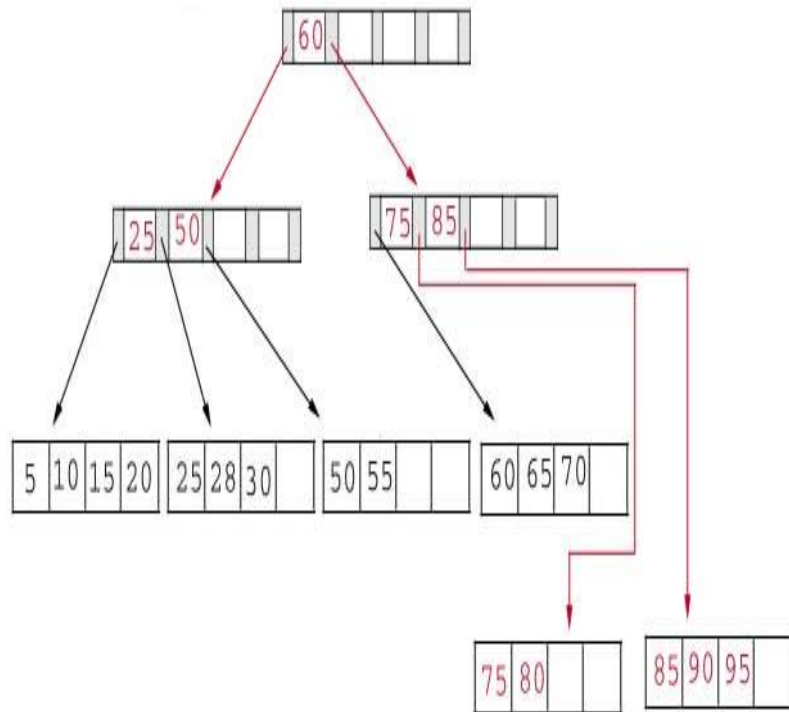
- ❑ As our example, we consider the B+ tree after we added 95 as a key.
- ❑ As a refresher this tree is printed in the following table.



Delete **70** from the B+ Tree

- ❑ We begin by deleting the record with key **70** from the B+ tree.
- ❑ This record is in a **leaf page** containing **60, 65 and 70**.
- ❑ This page will contain 2 records after the deletion.
- ❑ Since our fill factor is 50% or (2 records) we simply delete 70 from the leaf node.
- ❑ The following table shows the B+ tree after the deletion.



Delete Record with Key 70



Delete **25** from the B+ tree

- ❑ Delete the record containing **25** from the B+ tree.
- ❑ This record is found in the leaf node containing 25, 28, and 30.
- ❑ The fill factor will be 50% after the deletion; however, **25 appears in the index page.**
- ❑ Thus, when we delete 25 we must replace it with 28 in the index page.


Delete 25 from the B+ tree




Delete 60 from the B+ tree

- ❑ Delete 60 from the B+ tree. This deletion is interesting for several reasons:
 1. The **leaf page** containing 60 (60 65) will be **below the fill factor** after the deletion. Thus, we must **combine leaf pages**.
 2. With recombined pages, the **index page** will be reduced by one key. Hence, it will also **fall below the fill factor**. Thus, we must **combine index pages**.
 3. **Reduce index level**: 60 appears as the only key in the root index page. Obviously, it will be removed with the deletion.

The following table shows the B+ tree after the deletion of 60. Notice that the tree contains a single index page.



Delete Record with Key 60



Exercise

- ❑ Create an initial B+ tree with a B+-tree with $p = 3$ and $p_{leaf} = 2$
 - First Insert 8, 5,
 - Then insert in the following sequence
 - 1, 7, 3, 12, 9, 6
 - At each insertion observe, if it is increasing width or height of the tree
- ❑ Deletion Sequence: 5, 12, 9


Use this visualization tool to enhance your understanding:
<https://www.cs.usfca.edu/~galles/visualization/BPlusTree.html>

Hashing

- ❑ Provides *rapid, non-sequential, direct access* to records, by providing a *Search Key* value
- ❑ **Hashing:** A **key record** field is used to calculate the record address by subjecting it to some calculation; a process called *hashing*.
 - A hash function is **computed on some attribute** of each record.
 - The result of the function specifies in which block of the file the record should be placed

Hashing

- A hash structure (or table or file) is a *generalization* of the simpler notion of an ordinary array
 - In an array, an arbitrary position can be examined in $O(1)$
- A hash function h is used to map keys into a range of *bucket numbers*



Hash Indexes

- ❑ The **hash index** is a secondary structure to access the file by using **hashing** on a search key
- ❑ The index entries are of the type
 - $\langle K, Pr \rangle$ or $\langle K, P \rangle$,
 - where Pr is a pointer to the record containing the key, or P is a pointer to the block(bucket) containing the record for that key.

Hash-Based Indexing

- What indexing technique can we use to support *range searches* (e.g., “Find s_name where gpa \geq 3.0)?
 - Tree-Based Indexing
- What about *equality selections* (e.g., “Find s_name where sid = 102”)?
 - Tree-Based Indexing
 - Hash-Based Indexing (*cannot support range searches!*)

Downsides of Indexes

- 1) Extra space**
- 2) Index creation**
- 3) Index maintenance**

Picking which indexes to create

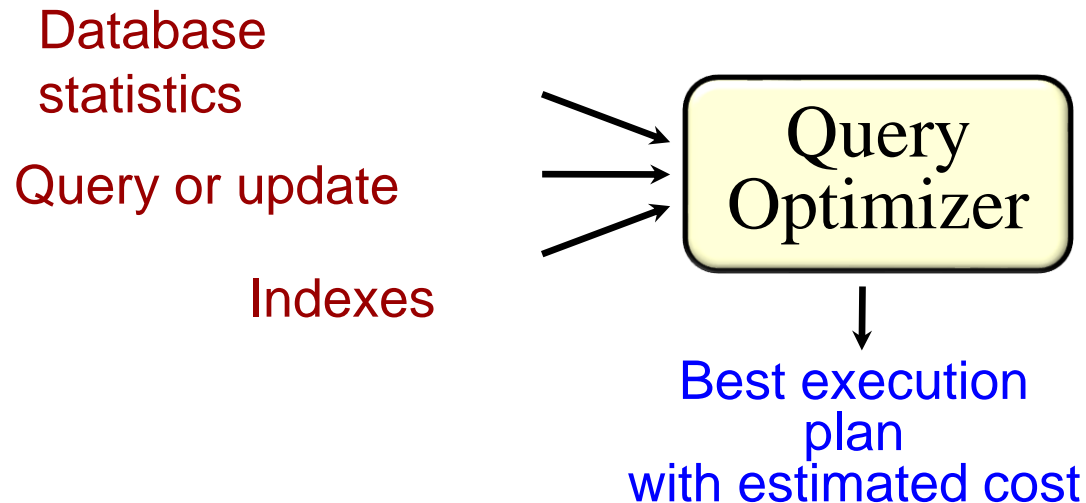
Benefit of an index depends on:

- Size of table (and possibly layout)
- Data distributions
- Query vs. update load

“Physical design advisors”

Input: database (statistics) and workload

Output: recommended indexes



Summary

- ❑ Indexes
- ❑ Multilevel Indexes
- ❑ Dynamic Multilevel Indexes Using B+-Trees
- ❑ Hash Indexes