

INTRODUCTION TO NEURAL NETWORKS

Lecture 10. Convolutional Neural Network (2)

Dr. Jingxin Liu

School of AI and Advanced Computing



Xi'an Jiaotong-Liverpool University

西交利物浦大學


Table of Contents

- I. Overfitting and Regularization
- II. Evaluation Metrics for Classification
- III. Summary
- IV. Future Works


Overfitting and Regularization

Overfitting: the model tries to learn too many details in the training data along with the noise from the training data. As a result, the model performance is very poor on unseen or test datasets.


Typical overfitting means that error on the training data is very low, but error on new instances is high.



Underfitted
(High bias error)



Good Fit/Robust
(Balance between
bias and variance)

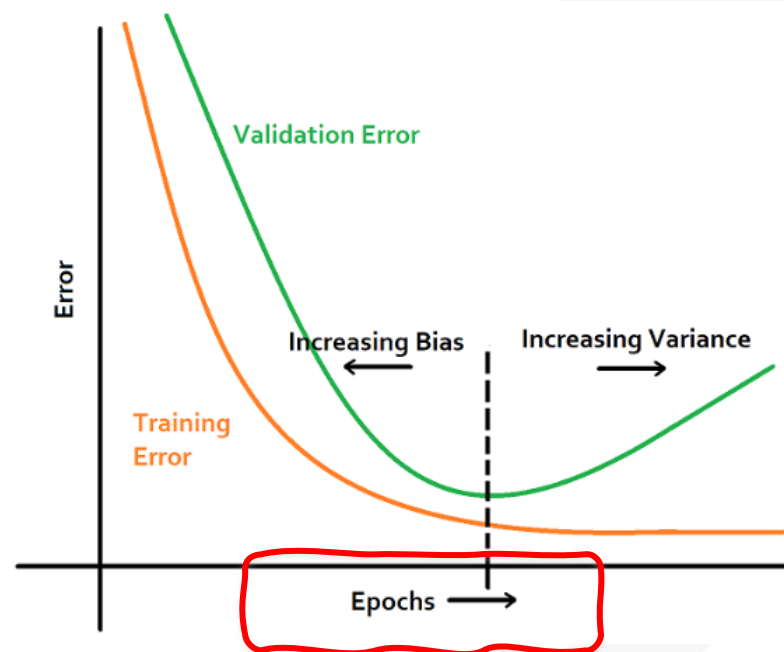
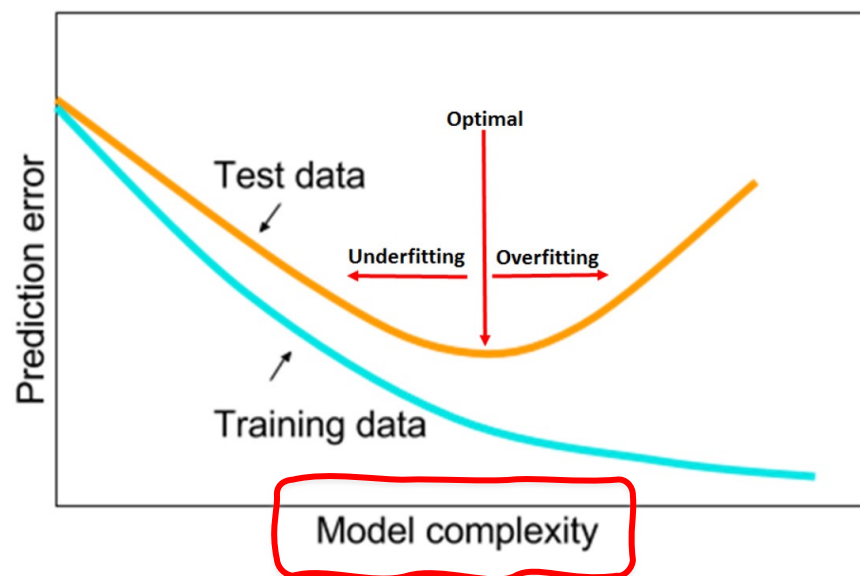


Overfitted
(High variance error)

Overfitting and Regularization

Overfitting is a general, **HUGELY IMPORTANT** problem for all machine learning algorithms.

- The size of the training dataset used is not enough.
- Data used for training is not cleaned and contains noise (garbage values) in it.
- The model is too complex.




Overfitting and Regularization

Regularization, in general: any method to **prevent overfitting** or help the optimization.


Regularization techniques:

- L1 regularization
- L2 regularization
-
- Simplifying the model
- Data augmentation
- Early stopping
- Dropouts
- Batch normalization


Early Stopping



Dropout



(a) Standard Neural Net




(b) After applying dropout.

```
vgg16.add(Flatten())
vgg16.add(Dense(4096))
vgg16.add(Activation('relu'))
vgg16.add(Dropout(0.5))
vgg16.add(Dense(4096))
vgg16.add(Activation('relu'))
vgg16.add(Dropout(0.5))
vgg16.add(Dense(n_classes))
vgg16.add(Activation('softmax'))
```

```
x = x.reshape(x.shape[0], -1)
x = F.relu(self.fc1(x))
x = F.dropout(x, 0.5) #dropout was included to combat overfitting
x = F.relu(self.fc2(x))
x = F.dropout(x, 0.5)
x = self.fc3(x)
```

Batch Normalization

Batch normalization (also known as **batch norm**) is a method used to make artificial neural networks faster and more stable through normalization of the layers' inputs by re-centering and re-scaling.



```
x = layers.Conv2D(filters1, (1, 1), strides=strides,
                  kernel_initializer='he_normal',
                  name=conv_name_base + '2a')(input_tensor)
x = layers.BatchNormalization(axis=bn_axis, name=bn_name_base + '2a')(x)
x = layers.Activation('relu')(x)

x = layers.Conv2D(filters2, kernel_size, padding='same',
                  kernel_initializer='he_normal',
                  name=conv_name_base + '2b')(x)
x = layers.BatchNormalization(axis=bn_axis, name=bn_name_base + '2b')(x)
x = layers.Activation('relu')(x)
```


Evaluation Metrics for Classification

Evaluation metrics are used to measure the quality of the statistical or machine learning model.

It is very important to **use multiple evaluation metrics to evaluate your model**. This is because a model may perform well using one measurement from one evaluation metric, but may perform poorly using another measurement from another evaluation metric. Using evaluation metrics are critical in ensuring that your model is operating correctly and optimally.

Evaluation metrics for **classification tasks**:

- Accuracy
- Confusion Matrix
- Precision and Recall
- F1 Score
- ROC curve
- Precision-Recall curve
-

Confusion Matrix

Confusion Matrix, also known as an **Error Matrix**, the Confusion Matrix is a two-dimensional matrix that allows visualization of the algorithm's performance. While this isn't an actual metric to use for evaluation, it's an important starting point.

In a binary classification, the matrix will be 2x2. If there are 3 classes, the matrix will be 3x3, and so on.

n=165	Predicted: NO	Predicted: YES
Actual: NO	50	10
Actual: YES	5	100

Actual	Elephant	25	3	0	2
	Monkey	3	53	2	3
	Fish	2	1	24	2
	Lion	1	0	2	71
		Elephant	Monkey	Fish	Lion
		Predicted			

Confusion Matrix

		Predicted Class		
		Positive	Negative	
Actual Class	Positive	True Positive (TP)	False Negative (FN) Type II Error	Sensitivity $\frac{TP}{(TP + FN)}$
	Negative	False Positive (FP) Type I Error	True Negative (TN)	Specificity $\frac{TN}{(TN + FP)}$
		Precision $\frac{TP}{(TP + FP)}$	Negative Predictive Value $\frac{TN}{(TN + FN)}$	Accuracy $\frac{TP + TN}{(TP + TN + FP + FN)}$

Accuracy

$$Accuracy = \frac{TP + TN}{TP + TN + FP + FN}$$

number of correctly classified examples
total number of examples

$$Precision = \frac{TP}{TP + FP}$$

$$Recall = \frac{TP}{TP + FN}$$

$$F1-score = \frac{2 \times Precision \times Recall}{Precision + Recall}$$

Precision and Recall

Define patients who have Covid-19 as **Positive** and those who do not have Covid-19 as **Negative**

- **Precision:** Out of the bunch of people diagnosed with Covid-19, how many actually have Covid-19?
- **Recall:** Out of a bunch of patients with Covid-19, how many people can be successfully detected Covid-19?

Spam filtering system, spam is **Positive**, normal mail is **Negative**

- **Precision:** How much of the blocked email is really spam?
- **Recall:** How much of all spam is blocked?

ROC Curve

An **ROC curve** (receiver operating characteristic curve) is a graph showing the performance of a classification model at all classification thresholds. This curve plots two parameters:



- True Positive Rate
- False Positive Rate

An ROC curve plots TPR vs. FPR at different classification thresholds. Lowering the classification threshold classifies more items as positive, thus increasing both False Positives and True Positives.

		Predicted Class		
		Positive	Negative	
Actual Class	Positive	True Positive (TP)	False Negative (FN) Type II Error	Sensitivity $\frac{TP}{(TP + FN)}$
	Negative	False Positive (FP) Type I Error	True Negative (TN)	Specificity $\frac{TN}{(TN + FP)}$
		Precision $\frac{TP}{(TP + FP)}$	Negative Predictive Value $\frac{TN}{(TN + FN)}$	Accuracy $\frac{TP + TN}{(TP + TN + FP + FN)}$



True Positive Rate

$$FPR = \frac{FP}{FP + TN}$$

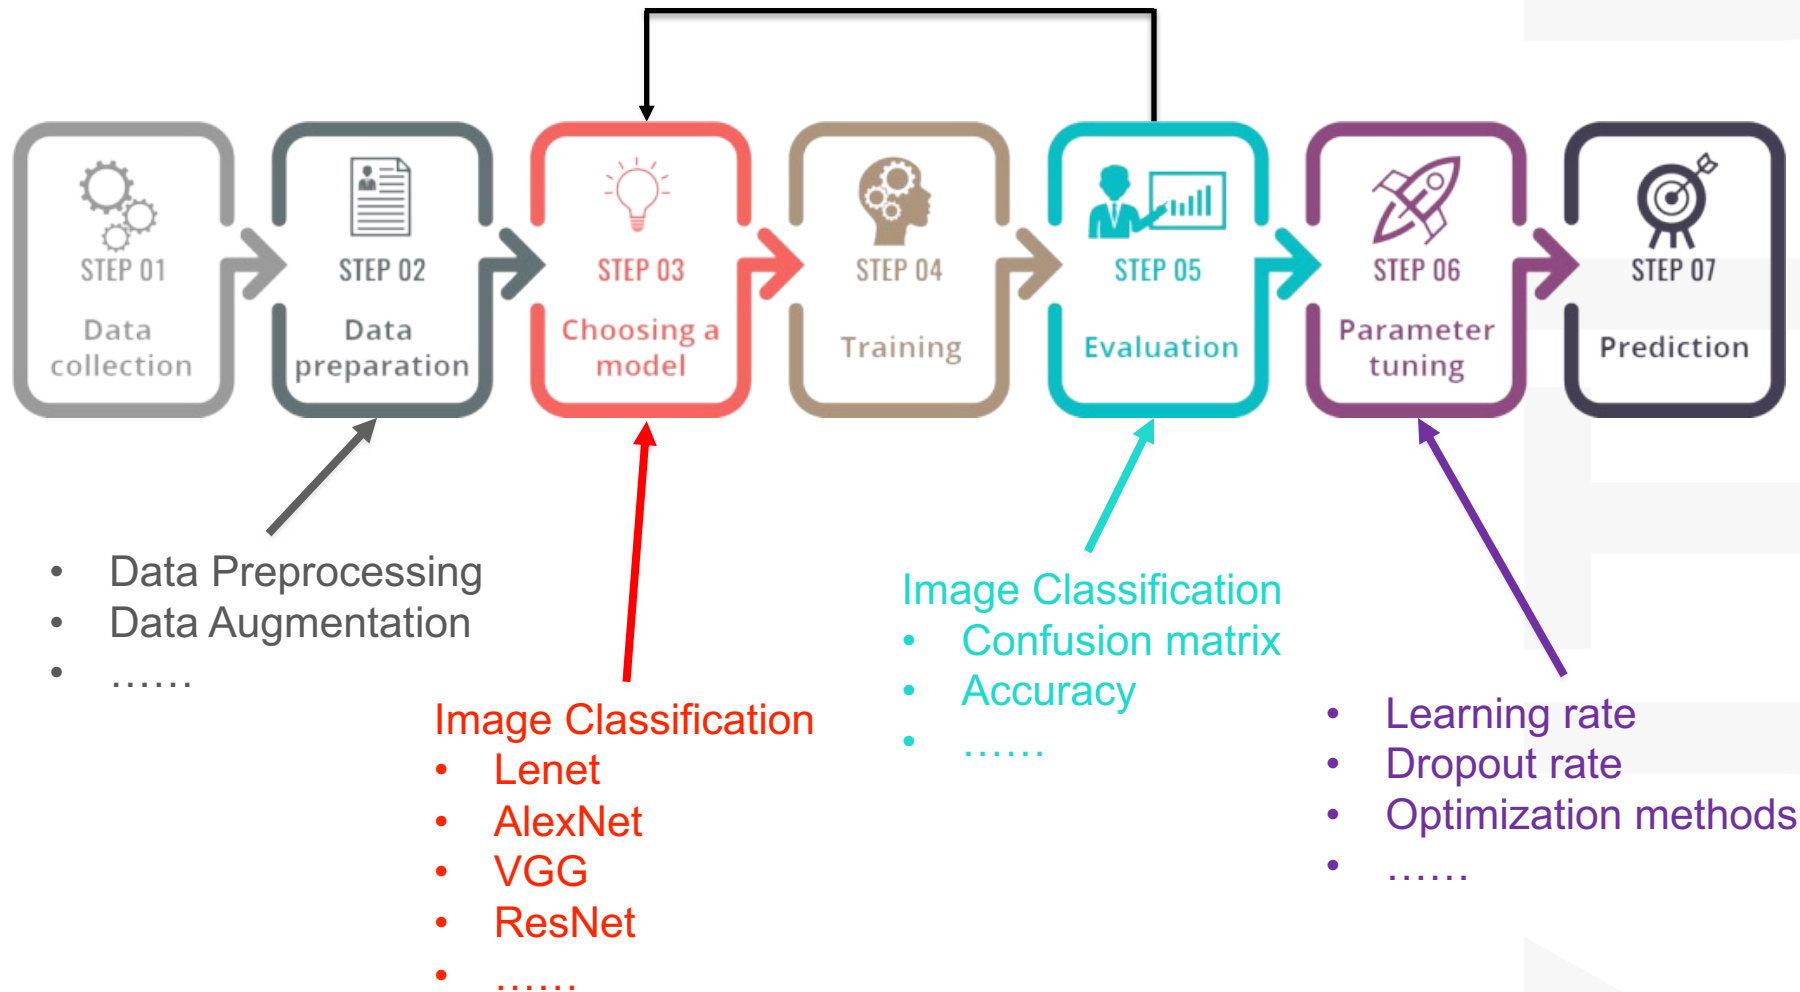


Precision-Recall Curve

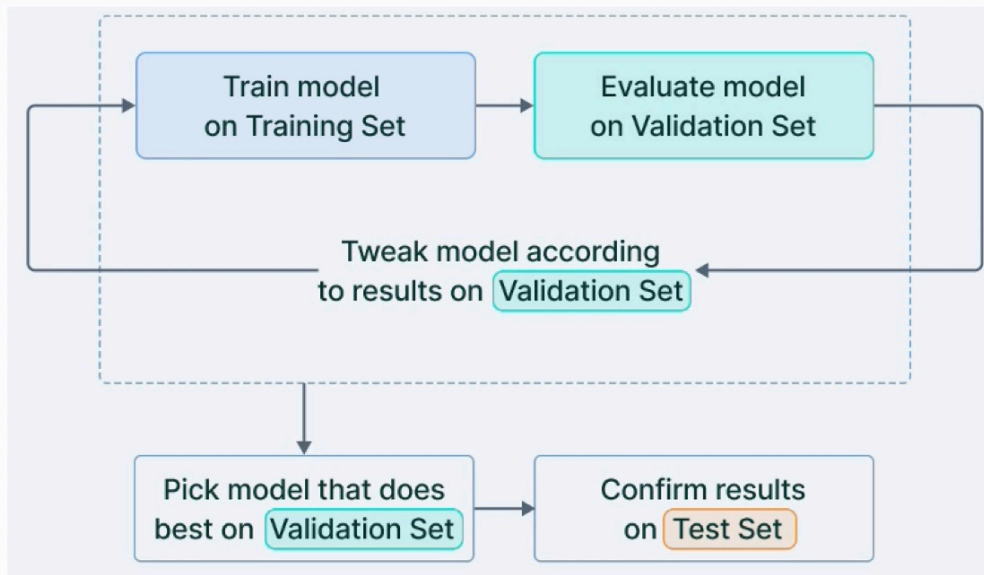
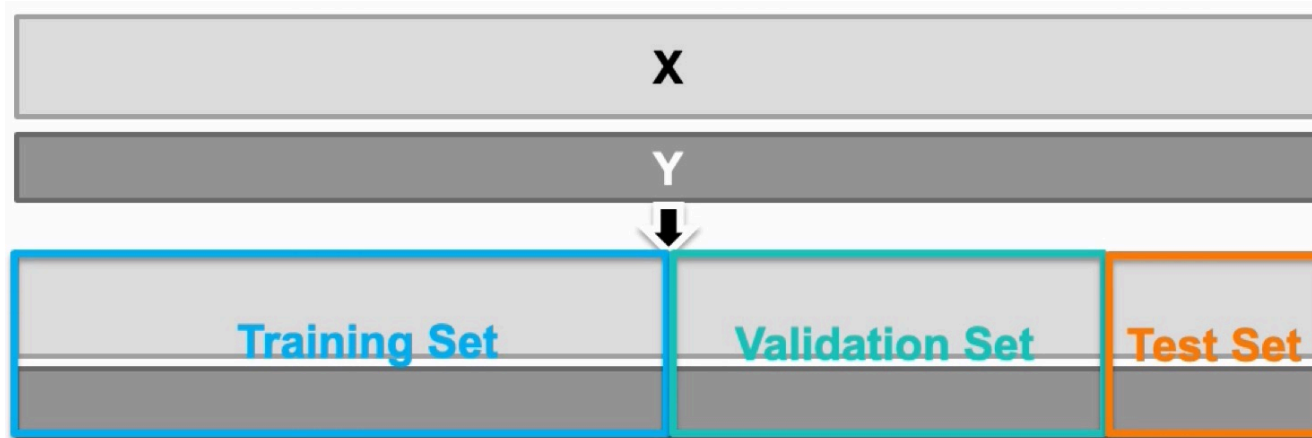
The **precision-recall curve** is used for evaluating the performance of binary classification algorithms. It is often used in situations where classes are **heavily imbalanced**. Also like ROC curves, precision-recall curves provide a graphical representation of a classifier's performance across many thresholds, rather than a single value.



Summary



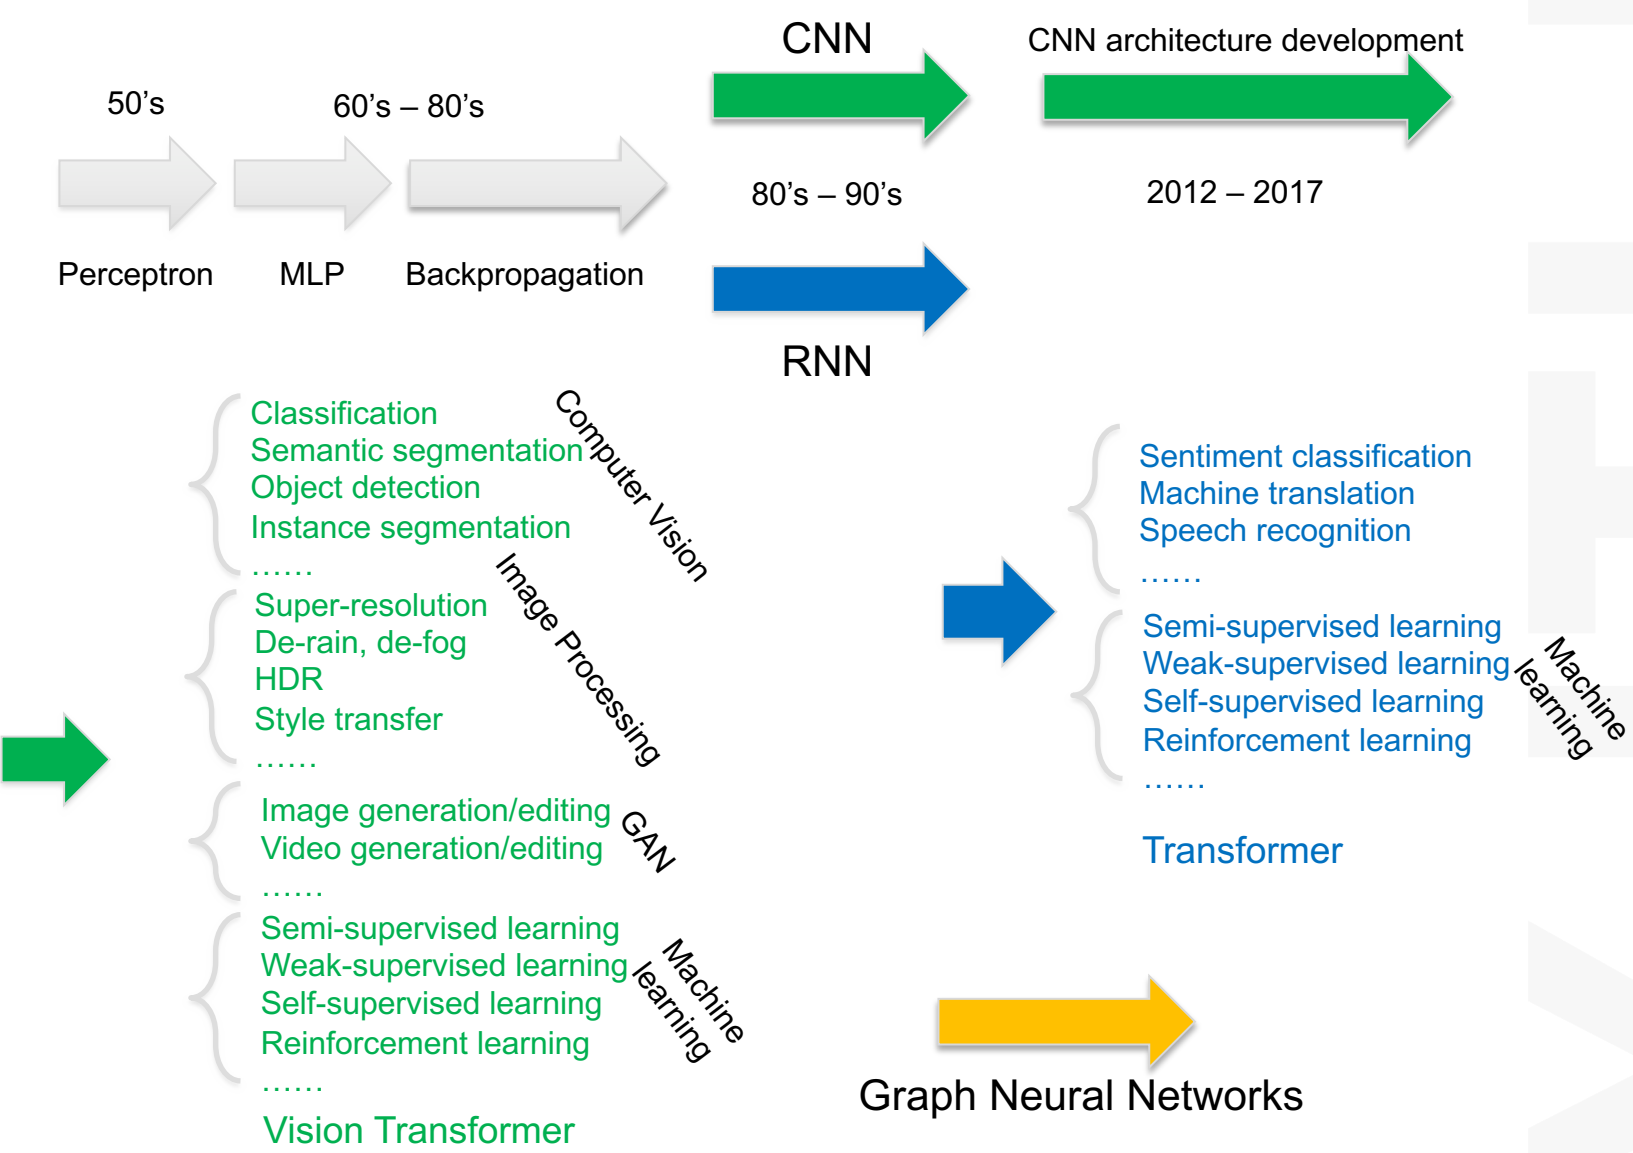
Summary



Summary

Lecture 1	Course Specification, AI, ML Introduction
Lecture 2	Linear Regression and Logistic Regression
Lecture 3	ANN Overview, Application, History, etc.
Lecture 4	Perceptron
Lecture 5	feed forward network, Forward Propagation
Lecture 6	Backpropagation
Lecture 7	Introduction to Recurrent Network
Lecture 8	Image Processing Fundamental
Lecture 9	Convolutional Neural Networks - 1
Lecture 10	Convolutional Neural Networks - 2

Future Works



Thank you !