

# DTS202TC Foundation of Parallel Computing

## Lecture 3

Hong-Bin Liu

Xi'an Jiaotong-Liverpool University

September 20, 2021

# Recall from last time

---



- Designing parallel program.
- Distributed-memory programming with MPI

# Some common problems



```
Sorry! You were supposed to get help about:  
    regex:invalid-name  
from the file:  
    help-regex.txt  
But I couldn't find that topic in the file. Sorry!
```

Chinese character in the Computer name

```
-----  
Sorry! You were supposed to get help about:  
    agent-not-found  
from the file:  
    help-plm-rsh.txt  
But I couldn't find that topic in the file. Sorry!  
-----
```

Missing openssh

# Some common problems



```
/cygdrive/d/xue/DTS202/F-hello
$ mpiexec -n 5 ./hello
-----
There are not enough slots available in the system to satisfy the 5
slots that were requested by the application:

./hello

Either request fewer slots for your application, or make more slots
available for use.

A "slot" is the Open MPI term for an allocatable unit where we can
launch a process. The number of slots available are defined by the
environment in which Open MPI processes are run:

1. Hostfile, via "slots=N" clauses (N defaults to number of
   processor cores if not provided)
2. The --host command line parameter, via a ":N" suffix on the
   hostname (N defaults to 1 if not provided)
3. Resource manager (e.g., SLURM, PBS/Torque, LSF, etc.)
4. If none of a hostfile, the --host command line parameter, or an
   RM is present, Open MPI defaults to the number of processor cores

In all the above cases, if you want Open MPI to default to the number
of hardware threads instead of the number of processor cores, use the
```

# Some common questions



outputted as `0`, even if it's just `NULL`.

4. Point-to-point communications are matched on the basis of tags and communicators. Collective communications don't use tags, so they're matched solely on the basis of the communicator and the *order* in which they're called. As an example, consider the calls to `MPI_Reduce` shown in Table 3.3. Suppose that each process calls `MPI_Reduce` with operator `MPI_SUM`, and destination process 0. At first glance, it might seem that after the two calls to `MPI_Reduce`, the value of `b` will be three, and the value of `d` will be six. However, the names of the memory locations are irrelevant to the matching, of the calls to `MPI_Reduce`. The *order* of the calls will determine the matching, so the value stored in `b` will be  $1 + 2 + 1 = 4$ , and the value stored in `d` will be  $2 + 1 + 2 = 5$ .

**Table 3.3** Multiple Calls to `MPI_Reduce`

Time	Process 0	Process 1	Process 2
0	<code>a = 1; c = 2</code>	<code>a = 1; c = 2</code>	<code>a = 1; c = 2</code>
1	<code>MPI_Reduce(&amp;a, &amp;b, ...)</code>	<code>MPI_Reduce(&amp;c, &amp;d, ...)</code>	<code>MPI_Reduce(&amp;a, &amp;b, ...)</code>
2	<code>MPI_Reduce(&amp;c, &amp;d, ...)</code>	<code>MPI_Reduce(&amp;a, &amp;b, ...)</code>	<code>MPI_Reduce(&amp;c, &amp;d, ...)</code>

- Assessment groups have been released.
- Decision is final.

[Download Excel Sheet](#)

# Goals for this week

---



Xi'an Jiaotong-Liverpool University  
西交利物浦大学

- Profiling tools
- Share-memory programming with Pthreads.

---

- Profiling Tools

- Pthreads

- Background
- Pthreads Basics
  - Our First Pthreads Program
  - Fundamental Concepts and Common Functions
- Matrix-Vector Multiplication
- Critical Sections and Synchronization
  - **Critical Sections**
  - Busy-Waiting
  - Mutex
  - Semaphore
  - Barrier and Condition Variables
- Thread Safety



# Profiling tools (profilers)

**Wikipedia: Profiling** is a form of [dynamic program analysis](#) that measures, for example, the space (memory) or time [complexity of a program](#), the [usage of particular instructions](#), or the frequency and duration of function calls. Most commonly, profiling information serves to aid [program optimization](#), and more specifically, [performance engineering](#).

- Gprof, Xcode Instruments
- mpiP
- HPCToolkit, caliper

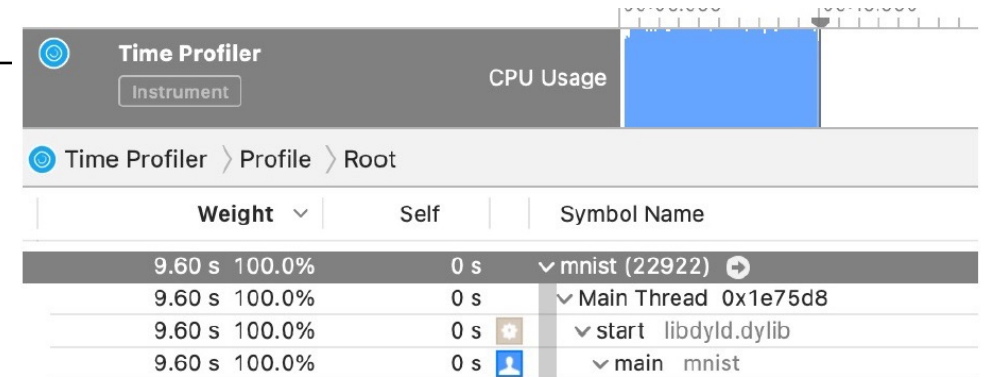


Figure 1: Instruments on Mac

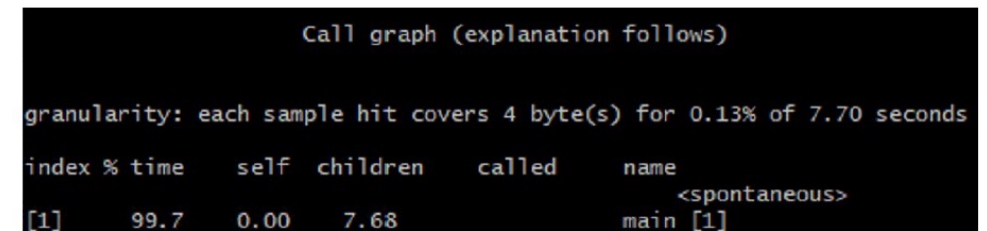


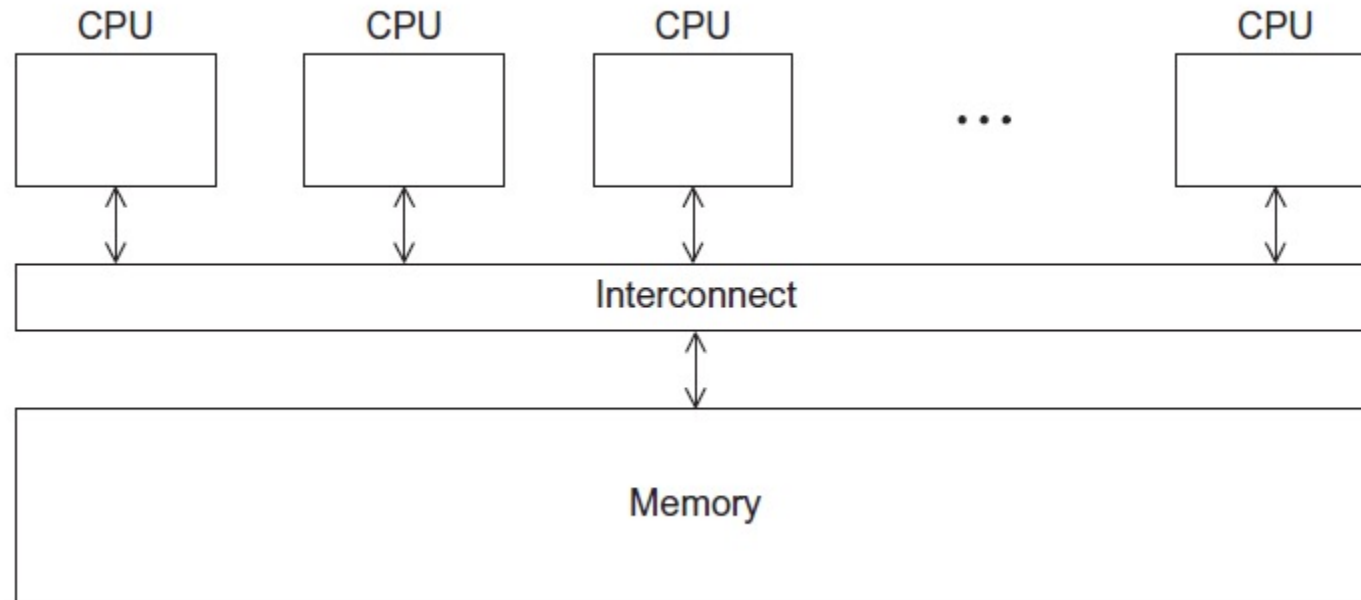
Figure 2: gprof on Cygwin Windows

---

<https://www.thegeekstuff.com/2012/08/gprof-tutorial/>

- 
- Background
  - Pthreads Basics
    - Our First Pthreads Program
    - Fundamental Concepts and Common Functions
  - Matrix-Vector Multiplication
  - Critical Sections and Synchronization
    - Critical Sections
    - Busy-Waiting
    - Mutex
    - Semaphore
    - Barrier and Condition Variables
  - Thread Safety

# A Shared Memory System



# Threads and Processes

---

- In shared-memory programming, a thread is an instance of a running (or suspended) program.
- Threads are analogous to a “light-weight” process.
- In a shared-memory program, a single process may have multiple threads of control.

- IEEE had a POSIX 1003 group that defined an interface to multithreaded programming
  - This is called Pthreads
  - A standard for Unix-like operating systems. The Pthreads API is only available on POSIX systems, however, we are using Cygwin on Windows.
  - A library that can be linked with C programs.
  - Provides primitives for thread management and synchronization
- Threads are peers, no parent/child relationship

- 
- Background
  - Pthreads Basics
    - Our First Pthreads Program
    - Fundamental Concepts and Common Functions
  - Matrix-Vector Multiplication
  - Critical Sections and Synchronization
    - Critical Sections
    - Busy-Waiting
    - Mutex
    - Semaphore
    - Barrier and Condition Variables
  - Thread Safety

# Hello World!



```
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>

/* Global variable: accessible to all threads */
int thread_count;

void *Hello(void* rank); /* Thread function */

int main(int argc, char* argv[]) {
    long thread; /* Use long in case of a 64-bit system */
    pthread_t* thread_handles;

    /* Get number of threads from command line */
    thread_count = strtol(argv[1], NULL, 10);

    thread_handles = malloc (thread_count*sizeof(pthread_t));
```

declares the various Pthreads functions, constants, types, etc.

global variable



# Hello World! Cont.



Create and  
start new  
threads

```
for (thread = 0; thread < thread_count; thread++)  
    pthread_create(&thread_handles[thread], NULL,  
                  Hello, (void*) thread);
```

Join threads

```
printf("Hello from the main thread\n");  
  
for (thread = 0; thread < thread_count; thread++)  
    pthread_join(thread_handles[thread], NULL);  
  
free(thread_handles);  
return 0;  
} /* main */
```

```
void *Hello(void* rank) {  
    long my_rank = (long) rank; /* Use long in case of 64-bit system */  
  
    printf("Hello from thread %ld of %d\n", my_rank, thread_count);  
  
    return NULL;  
} /* Hello */
```


# Compiling a Pthread program

---



```
gcc -g -Wall -o pth_hello pth_hello . c -lpthread
```

link in the Pthreads library



# Running a Pthreads program

---



`./pth_hello <number of threads>`

`./pth_hello 1`

Hello from the main thread  
Hello from thread 0 of 1

`./pth_hello 4`

Hello from the main thread  
Hello from thread 0 of 4  
Hello from thread 1 of 4  
Hello from thread 2 of 4  
Hello from thread 3 of 4

- `thread_count`
- Can introduce subtle and confusing bugs!
- Limit use of global variables to situations in which they're really needed.
  - Shared variables.


- 
- Background
  - Pthreads Basics
    - Our First Pthreads Program
    - Fundamental Concepts and Common Functions
  - Matrix-Vector Multiplication
  - Critical Sections and Synchronization
    - Critical Sections
    - Busy-Waiting
    - Mutex
    - Semaphore
    - Barrier and Condition Variables
    - Read-Write Locks
  - Thread Safety

- Processes in MPI are usually started by a script.
- In Pthreads, the threads are started by the program executable.
  - Include code to explicitly start the threads
  - Need data structures to store threads information

# Starting the Threads




pthread.h



pthread\_t

**One object  
for each  
thread.**



```
int pthread_create (  
    pthread_t*  thread_p /* out */,  
    const pthread_attr_t* attr_p /* in */,  
    void* (*start_routine) ( void ) /* in */,  
    void* arg_p /* in */ );
```

**The function creates and starts a new thread.**

- **Opaque**
- The actual data that they store is system-specific.
- Their data members aren't directly accessible to user code.
- However, the Pthreads standard guarantees that a pthread\_t object does store enough information to uniquely identify the thread with which it's associated.



# A Closer Look (1)



```
int pthread_create (  
    pthread_t* thread_p /* out */,  
    const pthread_attr_t* attr_p /* in */,  
    void* (*start_routine) ( void ) /* in */,  
    void* arg_p /* in */ );
```

We won't be using, so we just pass NULL.

Allocate before calling.

# A Closer Look (2)



```
int pthread_create (
    pthread_t* thread_p /* out */,
    const pthread_attr_t* attr_p /* in */,
    void* (*start_routine) ( void ) /* in */,
    void* arg_p /* in */ );
```

Pointer to the argument that should  
be passed to the function `start_routine`.

The function that the thread is to run.

# Function started by pthread\_create

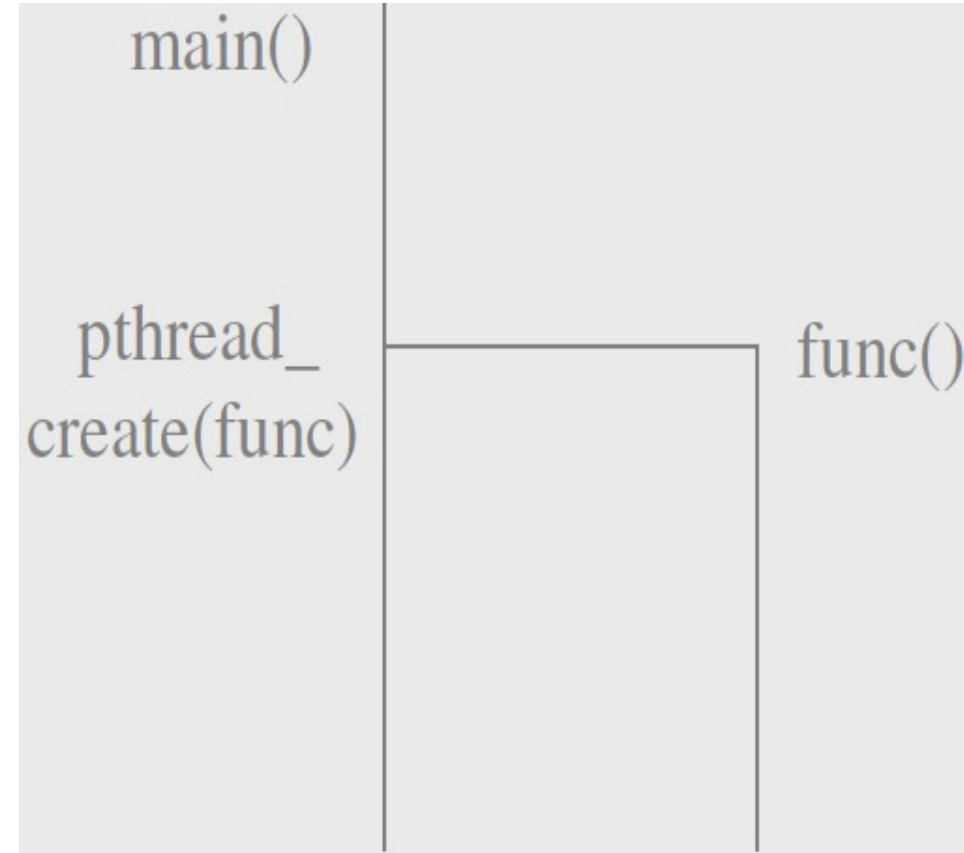


- Prototype:  
`void* thread_function ( void* args_p ) ;`
- Void\* can be cast to any pointer type in C.
- So args\_p can point to a list containing one or more values needed by thread\_function.
- Similarly, the return value of thread\_function can point to a list of one or more values.

# Example of Thread Creation



```
void *func(void *arg) {  
    int *i=arg;  
    .....  
}  
void main()  
{  
    int X; pthread_t id;  
    ....  
    pthread_create(&id, NULL,  
                  func, &X);  
    ...  
}
```



```
void pthread_exit(void *status)
```

- Terminates the **currently running** thread.

- We call the function `pthread_join` once for each thread.
- ```
int pthread_join(pthread_t new_id,      /* in */
                 void ** ret_val_p      /* out */ )
```

  - Waits for the thread with identifier `new_id` to terminate, either by returning or by calling `pthread_exit()`
- A single call to `pthread_join` will wait for the thread associated with the `pthread_t` object to complete.


# Thread Joining Example

---



```
void *func(void *) { ..... }  
pthread_t id; int X;  
pthread_create(&id, NULL, func, &X);  
.....  
pthread_join(id, NULL);  
.....
```

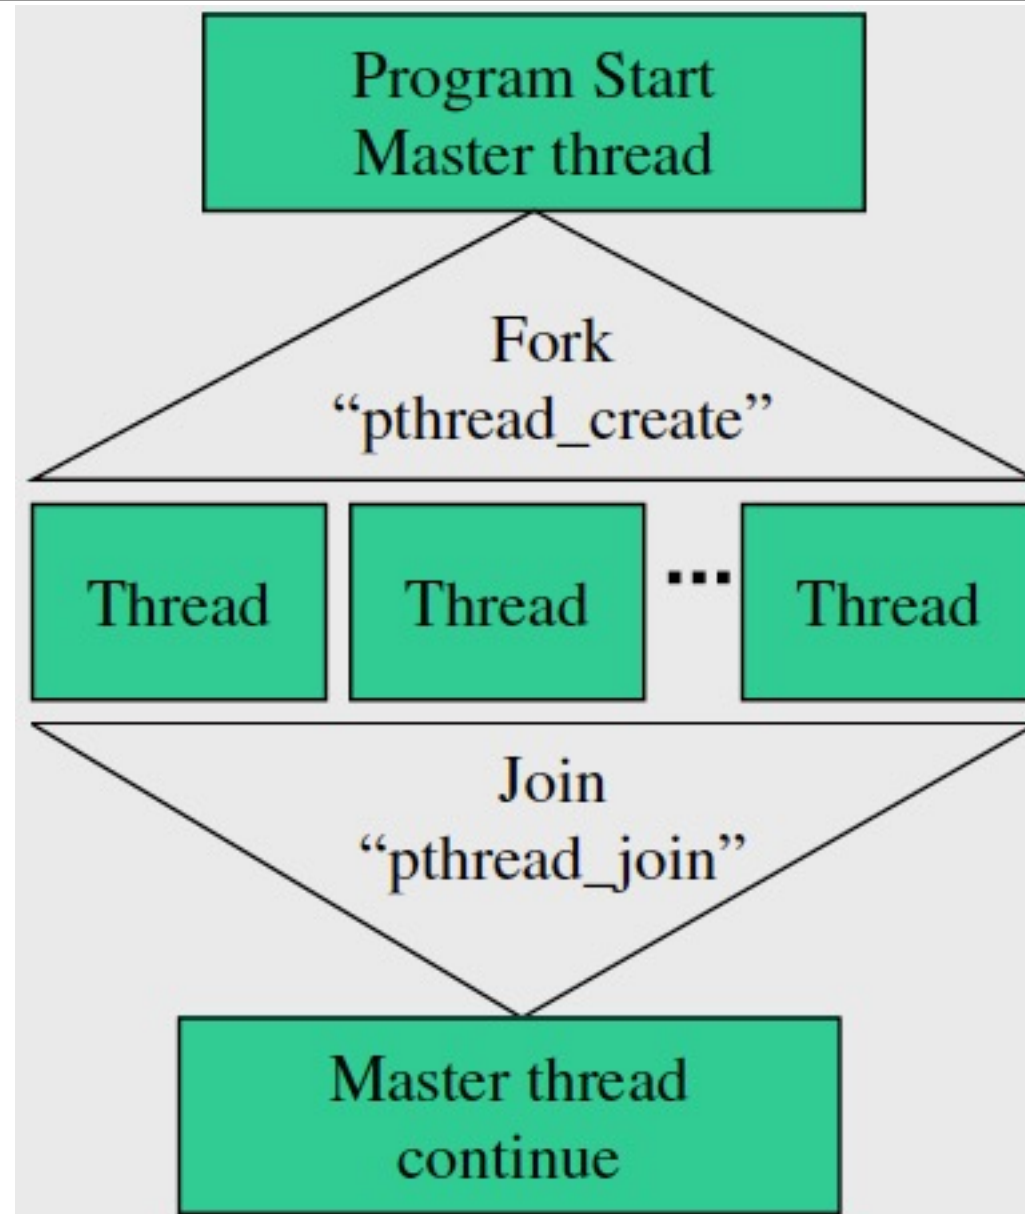
# Running the Threads



Main thread forks and joins two threads.



# Pthreads Programming Model




```
pthread_t pthread_self(void);
```

- To determine the thread ID of the calling thread

- 
- Background
  - Pthreads Basics
    - Our First Pthreads Program
    - Fundamental Concepts and Common Functions
  - Matrix-Vector Multiplication
  - Critical Sections and Synchronization
    - Critical Sections
    - Busy-Waiting
    - Mutex
    - Semaphore
    - Barrier and Condition Variables
  - Thread Safety

# Matrix-vector Multiplication



**FIGURE 3.11**

Matrix-vector multiplication

$$y_i = \sum_{j=0}^{n-1} a_{ij}x_j$$

# Serial Pseudo-code




```
/* For each row of A */
for (i = 0; i < m; i++) {
    y[i] = 0.0;
    /* For each element of the row and each element of x */
    for (j = 0; j < n; j++)
        y[i] += A[i][j]* x[j];
}
```

# Using Three Pthreads



| Thread | Components<br>of y |
|--------|--------------------|
| 0      | y[0], y[1]         |
| 1      | y[2], y[3]         |
| 2      | y[4], y[5]         |



thread 0

```
y[0] = 0.0;  
for (j = 0; j < n; j++)  
    y[0] += A[0][j]* x[j];
```



general case

```
y[i] = 0.0;  
for (j = 0; j < n; j++)  
    y[i] += A[i][j]*x[j];
```

# Pthreads Matrix-vector Multiplication



```
void *Pth_mat_vect(void* rank) {
    long my_rank = (long) rank;
    int i, j;
    int local_m = m/thread_count;
    int my_first_row = my_rank*local_m;
    int my_last_row = (my_rank+1)*local_m - 1;

    for (i = my_first_row; i <= my_last_row; i++) {
        y[i] = 0.0;
        for (j = 0; j < n; j++)
            y[i] += A[i][j]*x[j];
    }

    return NULL;
} /* Pth_mat_vect */
```

- 
- Background
  - Pthreads Basics
    - Our First Pthreads Program
    - Fundamental Concepts and Common Functions
  - Matrix-Vector Multiplication
  - Critical Sections and Synchronization
    - **Critical Sections**
    - Busy-Waiting
    - Mutex
    - Semaphore
    - Barrier and Condition Variables
  - Thread Safety



# Data Race in Pthreads Program



Xi'an Jiaotong-Liverpool University  
西交利物浦大学

```
static int s = 0;
```

Thread 1

```
for i = 0, n/2-1  
  s = s + f(A[i])
```

Thread 2

```
for i = n/2, n-1  
  s = s + f(A[i])
```

- Problem is a **race condition** on variable **s** in the program
- When multiple threads attempt to access a shared resource, at least one of the accesses is an update, and the accesses can result in an error, we have a **race condition**.

$$\pi = 4 \left( 1 - \frac{1}{3} + \frac{1}{5} - \frac{1}{7} + \cdots + (-1)^n \frac{1}{2n+1} + \cdots \right)$$

```
double factor = 1.0;
double sum = 0.0;
for (i = 0; i < n; i++, factor = -factor) {
    sum += factor/(2*i+1);
}
pi = 4.0*sum;
```

# A thread function for computing $\pi$



```
void* Thread_sum(void* rank) {
    long my_rank = (long) rank;
    double factor;
    long long i;
    long long my_n = n/thread_count;
    long long my_first_i = my_n*my_rank;
    long long my_last_i = my_first_i + my_n;

    if (my_first_i % 2 == 0) /* my_first_i is even */
        factor = 1.0;
    else /* my_first_i is odd */
        factor = -1.0;

    for (i = my_first_i; i < my_last_i; i++, factor = -factor) {
        sum += factor/(2*i+1);
    }

    return NULL;
} /* Thread_sum */
```

# Using a dual core processor



|           | $n$     |          |           |            |
|-----------|---------|----------|-----------|------------|
|           | $10^5$  | $10^6$   | $10^7$    | $10^8$     |
| $\pi$     | 3.14159 | 3.141593 | 3.1415927 | 3.14159265 |
| 1 Thread  | 3.14158 | 3.141592 | 3.1415926 | 3.14159264 |
| 2 Threads | 3.14158 | 3.141480 | 3.1413692 | 3.14164686 |

- Note that as we increase  $n$ , the estimate with one thread gets better and better.
- However, for larger values of  $n$ , the result computed by two threads actually gets worse.
- It matters if multiple threads try to update a single shared variable.

# Possible race condition



$y = \text{Compute}(\text{my rank});$   
 $x = x + y;$   
 $y$ : private,  $x$ : shared, initialized to 0

| Time | Thread 0                           | Thread 1                           |
|------|------------------------------------|------------------------------------|
| 1    | Started by main thread             |                                    |
| 2    | Call <code>Compute()</code>        | Started by main thread             |
| 3    | Assign $y = 1$                     | Call <code>Compute()</code>        |
| 4    | Put $x=0$ and $y=1$ into registers | Assign $y = 2$                     |
| 5    | Add 0 and 1                        | Put $x=0$ and $y=2$ into registers |
| 6    | Store 1 in memory location $x$     | Add 0 and 2                        |
| 7    |                                    | Store 2 in memory location $x$     |

A race condition or data race occurs when:

- Two processors (or two threads) access the same variable, and at least one does a write.
- The accesses are concurrent (not synchronized) so they could happen simultaneously.

- **Critical Section** is a block of code that updates a shared resources that can only be updated by one thread a time.

- 
- Background
  - Pthreads Basics
    - Our First Pthreads Program
    - Fundamental Concepts and Common Functions
  - Matrix-Vector Multiplication
  - Critical Sections and Synchronization
    - Critical Sections
    - Busy-Waiting
    - Mutex
    - Semaphore
    - Barrier and Condition Variables
    - Read-Write Locks
  - Thread Safety

- Create/exit/join
  - provide some form of synchronization,
  - at a very coarse level,
  - requires thread creation/destruction.
- Need for finer-grain synchronization
  - mutex locks
  - condition variables
  - ....
- PTHREADS provides a variety of synchronization facilities for threads to cooperate in accessing shared resources.



```
y = Compute(my_rank);  
while (flag != my_rank);  
x = x + y;  
flag++;
```

flag initialized to 0 by main thread

- Busy-waiting is not an ideal solution to the problem of controlling access to a critical section.
- Since thread 1 will execute the test over and over until thread 0 executes flag++ , if thread 0 is delayed, thread 1 will simply “spin” on the test, eating up CPU cycles. This can be positively disastrous for performance.

- A thread repeatedly tests a condition, but, effectively, does no useful work until the condition has the appropriate value.
- The busy-wait solution would work “provided the statements are executed exactly as they’re written.”
- Beware of optimizing compilers, though!

```
y = Compute(my_rank);  
while (flag != my_rank);  
x = x + y;  
flag++;
```

```
y = Compute(my_rank);  
x = x + y;  
while (flag != my_rank);  
flag++;
```

- Turning off compiler optimizations can seriously degrade performance.

# Pthreads Global Sum with Busy-waiting



```
void* Thread_sum(void* rank) {  
    long my_rank = (long) rank;  
    double factor;  
    long long i;  
    long long my_n = n/thread_count;  
    long long my_first_i = my_n*my_rank;  
    long long my_last_i = my_first_i + my_n;
```

```
    if (my_first_i % 2 == 0)  
        factor = 1.0;  
    else  
        factor = -1.0;
```

```
    for (i = my_first_i; i < my_last_i; i++, factor = -factor) {  
        while (flag != my_rank);  
        sum += factor/(2*i+1);  
        flag = (flag+1) % thread_count;  
    }
```

```
    return NULL;  
} /* Thread_sum */
```

When  $n=10^8$ , the serial sum is consistently faster than the parallel sum.

```
void* Thread_sum(void* rank) {
    long my_rank = (long) rank;
    double factor, my_sum = 0.0;
    long long i;
    long long my_n = n/thread_count;
    long long my_first_i = my_n*my_rank;
    long long my_last_i = my_first_i + my_n;

    if (my_first_i % 2 == 0)
        factor = 1.0;
    else
        factor = -1.0;

    for (i = my_first_i; i < my_last_i; i++, factor = -factor)
        my_sum += factor/(2*i+1);

    while (flag != my_rank);
    sum += my_sum;
    flag = (flag+1) % thread_count;

    return NULL;
} /* Thread_sum */
```

- Background
- Pthreads Basics
  - Our First Pthreads Program
  - Fundamental Concepts and Common Functions
- Matrix-Vector Multiplication
- Critical Sections and Synchronization
  - Critical Sections
  - Busy-Waiting
  - Mutex
  - Semaphore
  - Barrier and Condition Variables
  - Read-Write Locks
- Thread Safety

- A thread that is busy-waiting may continually use the CPU accomplishing nothing.
- Mutex (mutual exclusion) is a special type of **variable** that can be used to restrict access to a critical section to a single thread at a time.
- Used to guarantee that one thread “excludes” all other threads while it executes the critical section.

- The Pthreads standard includes a special type for mutexes: `pthread_mutex_t`. Need to be initialized by the system before it's used.

```
int pthread_mutex_init(  
    pthread_mutex_t*      mutex_p    /* out */  
    const pthread_mutexattr_t* attr_p /* in  */);
```

- Creates a new mutex lock.
- Attribute: normal, recursive, errorcheck  
(Ignore the second argument in this book, just pass NULL.)

- When a Pthreads program finishes using a mutex, it should call `pthread_mutex_destroy`

```
int pthread_mutex_destroy(pthread_mutex_t* mutex_p /* in/out */);
```

- In order to gain access to a critical section a thread calls `pthread_mutex_lock`

```
int pthread_mutex_lock(pthread_mutex_t* mutex_p /* in/out */);
```

- Tries to acquire the lock specified by mutex.
- If mutex is already locked, then calling thread blocks until mutex is unlocked.



- When a thread is finished executing the code in a critical section, it should call `pthread_mutex_unlock`

```
int pthread_mutex_unlock(pthread_mutex_t* mutex_p /* in/out */);
```

- If calling thread has mutex currently locked, this will unlock the mutex.
- If other threads are blocked waiting on this mutex, one will unblock and acquire mutex, which one is determined by the scheduler.

# Global Sum Function that Uses a Mutex



```
void* Thread_sum(void* rank) {
    long my_rank = (long) rank;
    double factor;
    long long i;
    long long my_n = n/thread_count;
    long long my_first_i = my_n*my_rank;
    long long my_last_i = my_first_i + my_n;
    double my_sum = 0.0;

    if (my_first_i % 2 == 0)
        factor = 1.0;
    else
        factor = -1.0;

    for (i = my_first_i; i < my_last_i; i++, factor = -factor) {
        my_sum += factor/(2*i+1);
    }
    pthread_mutex_lock(&mutex);
    sum += my_sum;
    pthread_mutex_unlock(&mutex);

    return NULL;
} /* Thread_sum */
```

- 
- Background
  - Pthreads Basics
    - Our First Pthreads Program
    - Fundamental Concepts and Common Functions
  - Matrix-Vector Multiplication
  - Critical Sections and Synchronization
    - Critical Sections
    - Busy-Waiting
    - Mutex
    - Semaphore
    - Barrier and Condition Variables
    - Read-Write Locks
  - Thread Safety

- 
- Busy-waiting enforces the order threads access a critical section.
  - Using mutexes, the order is left to chance and the system.
  - There are applications where we need to control the order threads access the critical section.

# Problems with a Mutex Solution



```
/* n and product_matrix are shared and initialized by the main  
thread */  
/* product_matrix is initialized to be the identity matrix */
```

```
void* Thread_work(void* rank) {  
    long my_rank = (long) rank;  
    matrix_t my_mat = Allocate_matrix(n);  
    Generate_matrix(my_mat);  
    pthread_mutex_lock(&mutex);  
    Multiply_matrix(product_mat, my_mat);  
    pthread_mutex_unlock(&mutex);  
    Free_matrix(&my_mat);  
    return NULL;  
} /* Thread_work */
```

# A First Attempt at Sending Messages Using Pthreads



```
/* messages has type char**. It's allocated in main. */
/* Each entry is set to NULL in main. */
void *Send_msg(void* rank) {
    long my_rank = (long) rank;
    long dest = (my_rank + 1) % thread_count;
    long source = (my_rank + thread_count - 1) % thread_count;
    char* my_msg = malloc(MSG_MAX*sizeof(char));

    sprintf(my_msg, "Hello to %ld from %ld", dest, my_rank);
    messages[dest] = my_msg;

    if (messages[my_rank] != NULL) {
        printf("Thread %ld > %s\n", my_rank, messages[my_rank]);
    } else {
        printf("Thread %ld > No message from %ld\n", my_rank, source);
    }

    return NULL;
} /* Send_msg */
```

No receiving message for  
multithreads

# Possible Solutions with Problems



**S1:** Change **if** statement to **while** statement:

```
-if (messages[my_rank] != NULL)
    printf("Thread %ld > %s\n", my_rank, messages[my_rank]);
-while (messages[my_rank] = NULL);
    printf("Thread %ld > %s\n", my_rank, messages[my_rank]);
•However, there is a busy-waiting problem.
```

•**S2: Mutex Lock:**

```
...
Pthread_mutex_lock(mutex[dest]);
...
Messages[dest]=msg;
Pthread_mutex_unlock(mutex[dest]);
...
Pthread_mutex_lock(mutex[myrank]);
printf("Thread %ld > %s\n", my_rank, messages[my_rank]);
...
```

•However, it still will crash.



- Semaphores can be thought of as a special type of `unsigned int`, so they can take on the values 0, 1, 2, . . . . When they take on the values 0 and 1, called a binary semaphore.
- For `binary semaphore`, 0 corresponds to a locked mutex, and 1 corresponds to an unlocked mutex.
  - Initialized to 1(unlocked).
  - Before the critical, call `sem_wait`. Executing `sem_wait` will block if the semaphore is 0. If the semaphore is nonzero, it will decrement the semaphore and proceed.
  - After executing the critical, call `sem_post`, which increments the semaphore, and a thread waiting in `sem_wait` can proceed.



```
#include <semaphore.h>
```

Semaphores are not part of  
Pthreads; you need to add this.

```
int sem_init(  
    sem_t*      semaphore_p    /* out */,  
    int         shared          /* in  */,  
    unsigned    initial_val     /* in  */);
```

Initialize the semaphore descriptor.

Unlock a semaphore.

```
int sem_destroy(sem_t* semaphore_p /* in/out */);  
int sem_post(sem_t* semaphore_p /* in/out */);  
int sem_wait(sem_t* semaphore_p /* in/out */);
```

Destroy a semaphore.

Lock a semaphore.

# Send Messages Using Semaphores



```
1  /* messages is allocated and initialized to NULL in main */
2  /* semaphores is allocated and initialized to 0 (locked) in
   main */
3  void* Send_msg(void* rank) {
4      long my_rank = (long) rank;
5      long dest = (my_rank + 1) % thread_count;
6      char* my_msg = malloc(MSG_MAX*sizeof(char));
7
8      sprintf(my_msg, "Hello to %ld from %ld", dest, my_rank);
9      messages[dest] = my_msg;
10     sem_post(&semaphores[dest])
11         /* "Unlock" the semaphore of dest */
12
13     /* Wait for our semaphore to be unlocked */
14     sem_wait(&semaphores[my_rank]);
15     printf("Thread %ld > %s\n", my_rank, messages[my_rank]);
16
17     return NULL;
18 }
```

- The message-sending problem didn't involve a critical section.
- Thread my\_rank couldn't proceed until thread source had finished creating the message.
- When a thread can't proceed until another thread has taken some action, is sometimes called **producer-consumer synchronization**.

- 
- Background
  - Pthreads Basics
    - Our First Pthreads Program
    - Fundamental Concepts and Common Functions
  - Matrix-Vector Multiplication
  - Critical Sections and Synchronization
    - Critical Sections
    - Busy-Waiting
    - Mutex
    - Semaphore
    - Barrier and Condition Variables
  - Thread Safety

- 
- Synchronizing the threads to make sure that they all are at the same point in a program is called a **barrier**.
  - No thread can cross the barrier until all the threads have reached it.

# Using Barriers to Time the Slowest Thread



```
/* Shared */
double elapsed_time;
. . .
/* Private */
double my_start, my_finish, my_elapsed;
. . .
Synchronize threads;
Store current time in my_start;
/* Execute timed code */
. . .
Store current time in my_finish;
my_elapsed = my_finish - my_start;

elapsed = Maximum of my_elapsed values;
```

# Using Barriers for Debugging

---



```
point in program we want to reach;  
barrier;  
if (my_rank == 0) {  
    printf("All threads reached this point\n");  
    fflush(stdout);  
}
```

- Implementing a barrier using busy-waiting and a mutex is straightforward.
- We use a shared **counter** protected by the **mutex**.
- When the counter indicates that every thread has entered the critical section, threads can leave the **busy-wait** loop.

# Busy-waiting and a Mutex



```
/* Shared and initialized by the main thread */
int counter; /* Initialize to 0 */
int thread_count;
pthread_mutex_t barrier_mutex;
. . .

void* Thread_work(. . .) {
    . . .
    /* Barrier */
    pthread_mutex_lock(&barrier_mutex);
    counter++;
    pthread_mutex_unlock(&barrier_mutex);
    while (counter < thread_count);
    . . .
}
```

We need one counter variable for each instance of the barrier, otherwise problems are likely to occur.



# Implementing a Barrier with Semaphores



```
/* Shared variables */
int counter; /* Initialize to 0 */
sem_t count_sem; /* Initialize to 1 */
sem_t barrier_sem; /* Initialize to 0 */
. . .
void* Thread_work(...) {
    . . .
    /* Barrier */
    sem_wait(&count_sem);
    if (counter == thread_count - 1) {
        counter = 0;
        sem_post(&count_sem);
        for (j = 0; j < thread_count - 1; j++)
            sem_post(&barrier_sem);
    } else {
        counter++;
        sem_post(&count_sem);
        sem_wait(&barrier_sem);
    }
    . . .
}
```

protects the counter

block threads that have entered the barrier

Reusing  
**barrier\_sem**  
results in a  
race condition.

- A **condition variable** is a data object that allows a thread to suspend execution until a certain event or condition occurs.
- When the event or condition occurs, another thread can **signal** the thread to “wake up.”
- A **condition variable** is always associated with a **mutex**.

```
lock mutex;  
if condition has occurred  
    signal thread(s);  
else {  
    unlock the mutex and block;  
    /* when thread is unblocked, mutex is relocked */  
}  
unlock mutex;
```

- Condition variables in Pthreads have type `pthread_cond_t`.

```
int pthread_cond_init(  
    pthread_cond_t*      cond_p      /* out */,  
    const pthread_condattr_t* cond_attr_p /* in */);
```

- Creates a new condition variable `cond`.
- Attribute: ignore for now.

```
int pthread_cond_destroy(pthread_cond_t* cond_p /* in/out */);
```

- Destroys the condition variable `cond`.

```
int pthread_cond_signal(pthread_cond_t* cond_var_p /* in/out */);
```


- Unblock one of the blocked threads waiting on cond.
- Which one is determined by scheduler.
- If no thread waiting, then signal is a no-op.

```
int pthread_cond_broadcast(pthread_cond_t* cond_var_p /* in/out */);
```

- Unblock all of the blocked threads waiting on cond.
- If no thread waiting, then signal is a no-op.

```
int pthread_cond_wait(  
    pthread_cond_t*    cond_var_p    /* in/out */,  
    pthread_mutex_t*   mutex_p       /* in/out */);
```

- Blocks the calling thread, waiting on cond.
- Unlocks the mutex.



```
pthread_mutex_unlock(&mutex_p);  
wait_on_signal(&cond_var_p);  
pthread_mutex_lock(&mutex_p);
```

```
/* Shared */
int counter = 0;
pthread_mutex_t mutex;
pthread_cond_t cond_var;
. . .
void* Thread_work(. . .) {
    . . .
    /* Barrier */
    pthread_mutex_lock(&mutex);
    counter++;
    if (counter == thread_count) {
        counter = 0;
        pthread_cond_broadcast(&cond_var);
    } else {
        while (pthread_cond_wait(&cond_var, &mutex) != 0);
    }
    pthread_mutex_unlock(&mutex);
    . . .
}
```

- 
- Background
  - Pthreads Basics
    - Our First Pthreads Program
    - Fundamental Concepts and Common Functions
  - Matrix-Vector Multiplication
  - Critical Sections and Synchronization
    - Critical Sections
    - Busy-Waiting
    - Mutex
    - Semaphore
    - Barrier and Condition Variables
    - Read-Write Locks
  - Thread Safety



- A block of code is **thread-safe** if it can be simultaneously executed by multiple threads without causing problems.

- 
- Suppose we want to use multiple threads to “tokenize” a file that consists of ordinary English text.
  - The tokens are just contiguous sequences of characters separated from the rest of the text by white-space — a space, a tab, or a newline.

- Divide the input file into lines of text and assign the lines to the threads in a round-robin fashion.
- The first line goes to thread 0, the second goes to thread 1, . . . , the  $t$ th goes to thread  $t$ , the  $t + 1$ st goes to thread 0, etc.
- We can serialize access to the lines of input using semaphores.
- After a thread has read a single line of input, it can tokenize the line using the `strtok` function.

# The Strtok Function



```
char* strtok(  
    char*      string      /* in/out */,  
    const char* separators /* in    */);
```

- The **idea** is that in the first call, **strtok** caches a pointer to string, and for subsequent calls it returns successive tokens **taken from the cached copy**.
- The first time it's called the string argument should be the text to be tokenized.(One line of input.)
- For subsequent calls, the first argument should be NULL.

# Multi-threaded Tokenizer



```
void* Tokenize(void* rank) {
    long my_rank = (long) rank;
    int count;
    int next = (my_rank + 1) % thread_count;
    char *fg_rv;
    char my_line[MAX];
    char *my_string;

    sem_wait(&sems[my_rank]);
    fg_rv = fgets(my_line, MAX, stdin);
    sem_post(&sems[next]);
    while (fg_rv != NULL) {
        printf("Thread %ld > my line = %s", my_rank, my_line);

        count = 0;
        my_string = strtok(my_line, " \t\n");
        while ( my_string != NULL ) {
            count++;
            printf("Thread %ld > string %d = %s\n", my_rank, count,
                my_string);
            my_string = strtok(NULL, " \t\n");
        }

        sem_wait(&sems[my_rank]);
        fg_rv = fgets(my_line, MAX, stdin);
        sem_post(&sems[next]);
    }

    return NULL;
} /* Tokenize */
```

- It correctly tokenizes the input stream.

Pease porridge hot.  
Pease porridge cold.  
Pease porridge in the pot  
Nine days old.

# Running with Two Threads



```
Thread 0 > my line = Pease porridge hot.  
Thread 0 > string 1 = Pease  
Thread 0 > string 2 = porridge  
Thread 0 > string 3 = hot.  
Thread 1 > my line = Pease porridge cold.  
Thread 0 > my line = Pease porridge in the pot  
Thread 0 > string 1 = Pease  
Thread 0 > string 2 = porridge  
Thread 0 > string 3 = in  
Thread 0 > string 4 = the  
Thread 0 > string 5 = pot  
Thread 1 > string 1 = Pease  
Thread 1 > my line = Nine days old.  
Thread 1 > string 1 = Nine  
Thread 1 > string 2 = days  
Thread 1 > string 3 = old.
```

Oops!



# What happened?



- `strtok` caches the input line by declaring a variable to have static storage class.
- This causes the value stored in this variable to persist from one call to the next.
- Unfortunately for us, `this cached string is shared, not private`.
- Thus, thread 0's call to `strtok` with the third line of the input has apparently `overwritten` the contents of thread 1's call with the second line.
- So the `strtok` function is not thread-safe. If multiple threads call it simultaneously, the output may not be correct.



- Regrettably, it's not uncommon for C library functions to fail to be thread-safe.
- The random number generator `random` in `stdlib.h`.
- The time conversion function `localtime` in `time.h`.

# “re-entrant” (Thread Safe) Functions



- In some cases, the C standard specifies an alternate, thread-safe, version of a function.

```
char* strtok_r(  
    char*      string      /* in/out */,  
    const char* separators, /* in    */,  
    char**     saveptr_p   /* in/out */);
```

---

Incorrect programs can produce  
correct output!

# Summary

---



Xi'an Jiaotong-Liverpool University  
西交利物浦大学